

Literature Survey

Shubhra Prakash 23BIT0087

1. Digital twin-driven deep reinforcement learning for real-time optimisation in dynamic AGV systems

This paper explores how to merge digital twins (DT) with deep reinforcement learning (DRL) to manage automated guided vehicles in dynamic manufacturing environments. Instead of just using the digital twin as a sandbox to test things after the fact, the system actively uses it to train the AI and make routing decisions on the fly.

Strengths	Weaknesses
The architecture is incredibly adaptable to sudden environmental changes because the DT provides a continuous, real-time feedback loop to the AI. And the system successfully learns to navigate without any collisions by heavily penalizing unsafe paths during its training phase.	The setup relies on a centralized decision-making structure, which could lead to computational bottlenecks and reduced responsiveness if scaled up to a massive multi-agent environment. But building and maintaining these high-fidelity, real-time synchronized digital twins is very expensive and resource-intensive

2. Disaster Management in the Era of Agentic AI Systems: A Vision for Collective Human-Machine Intelligence

This paper introduces "Disaster Copilot," which is exactly the kind of system we are trying to build. It uses a central orchestrator like a multi-modality large language model to manage specialized sub-agents that handle tasks like early warning, risk analytics, and resource allocation during disasters.

Strengths	Weaknesses
They manage to brilliantly solves the problem of data silos by keeping all the different specialized AI models in one cohesive, federated ecosystem. It features a distributed memory bank that retains past disaster knowledge, which acts as a permanent institutional memory even if the human personnel change.	The specialized sub-agents might become brittle and struggle to adapt when facing completely new environmental conditions that fall outside of their training data. There is currently no universally adopted communication protocol, which remains a structural barrier for making different decentralized systems interoperate seamlessly.

3. Intelligent multi-agent digital twin architecture for modelling and monitoring catastrophic events

This paper proposes a formally grounded framework that uses multi-agent digital twins specifically to monitor and manage catastrophic events, like floods, in cyber-physical environments. It maps out how agents communicate and make decisions using mathematical stochastic spatial-temporal models and timed Petri nets.

Strengths	Weaknesses:
It provides a highly structured and rigorous mathematical foundation for modelling how disaster processes unfold over time and space.	Because the paper focuses so heavily on the theoretical architecture, it still requires future experimental validation and real-world deployment to prove its practical effectiveness.
The framework is perfectly suited for distributed monitoring, scalable coordination, and adaptive risk assessment in real-time.	

4. Enhancing Digital Twins with Deep Reinforcement Learning: A Use Case in Maintenance Prioritization

This study looks at using DRL inside a digital twin to figure out which machines need maintenance first on an assembly line. It builds a closed-loop system where real-time physical data feeds into the AI to predict issues and prevent bottlenecks before they happen.

Strengths	Weaknesses
It proves that merging AI with digital twins can drastically outperform standard empirical or random decision-making.	The simulation relied on simplified assumptions—like assuming machine availability was constant—which doesn't reflect the messy, variable reality of the real world.
The reward functions are cleverly designed around actual production costs and sustainability, making the AI's decisions highly practical.	
	The specific configurations used in their lab might not easily generalize to other, more complex industrial setups without significant modifications.

5. Reinforcement learning and digital twin-driven optimization of production scheduling with the digital model playground

This paper introduces the "Digital Model Playground" (DMPG), an open-source Python framework designed to run discrete event simulations and train RL agents. It is built to be a flexible, user-friendly tool for creating digital shadows and twins for production scheduling.

Strengths	Weaknesses
It is completely open-source, adaptable, and performs just as fast as expensive commercial simulation software.	It hogs a massive amount of RAM during large-scale simulations because it doesn't delete data until all calculations are finished.
It can distribute heavy simulations to high-performance computing clusters and even output cool 3D visualizations using Unreal Engine.	Users need strong Python programming skills to use it, and the framework currently lacks built-in logic for simulating things like transport vehicles or complex logistics.

Conclusion

When we look at all five of these papers together, a really exciting picture emerges for our cloud-based disaster response system! The big theme connecting everything is the incredible potential of combining Digital Twins which are basically live, virtual copies of the real world with smart, learning AI agents to make split-second decisions in chaotic environments.

Instead of just using AI as a static tool, these studies show how we can build active ecosystems where multiple AI *specialists* talk to each other and learn on the fly. For example, the research highlights how we can use a main *orchestrator* AI to coordinate various sub-agents during a crisis, breaking down data silos so that rescue teams and logistics planners get exactly the insights they need. We also have clear, structured blueprints for how to make these distributed agents monitor catastrophic events and communicate seamlessly across a network.

At the same time, we see massive proof that plugging Reinforcement Learning into these digital twins allows the system to constantly adapt to its environment. This kind of setup can actively figure out safe routes and coordinate movements while completely avoiding collisions. It can also predict exactly what needs attention first by looking at real-time data and system constraints, which is perfect for prioritizing our scarce resources during an emergency.

What is really inspiring is that we do not even have to start from scratch to test these ideas. There are open-source simulation frameworks out there designed specifically to handle

these digital models, and they allow us to distribute the heavy computational work across cloud clusters to keep everything running fast.

Overall, by taking the real-time routing and predictive ideas from the factory-focused papers and merging them with the disaster-specific multi-agent architectures, we have a complete and practical blueprint. We can build a coordination system that doesn't just react to emergencies, but actively anticipates them.